

Final Project – Fall 2024

**AI and Machine Learning Hardware: Exploring
Specialized Architectures**

Students:

Kamar Modar Bdair – S23108164

Jwdee Fras Alamam – S23108505

Sedra Ayman Al-Jundi – S23108200

Instructor:

Dr. Mohmmad Amjad

Date:

14/11/2024

Table of Contents

Introduction	3
Main Body	3
-The Unique Demands of AI and ML Workloads	3
-Characteristics of AI and ML Workloads	5

Final Project – Fall 2024

-Tensor Processing Units (TPUs)	5
-Key Features	6
-Generations of TPUs	6
-Uses Cases	7
-Neuromorphic Chips	7
-Key Features	8
-Examples of Neuromorphic Chips	8
-Use Cases	9
-Comparison of TPUs and Neuromorphic Chips	9
-Impacts on Computer Architecture	9
Conclusion	10
References	12

Introduction

This decade is witnessing the phenomenal growth of artificial intelligence (AI) and machine learning (ML) technology. These technologies have changed the scene completely, from self-driving cars and natural language processing (NLP) to healthcare analytics and diagnostics. All these domains need effective algorithms for dealing with large amounts of data and performing complicated calculations. Computers would be useless, without hardware such as Central Processing Units (CPUs) and Graphics Processing Units (GPUs), which played a huge role in the development of AI, but the AI workloads of today are not suited for such architectures.

Final Project – Fall 2024

To meet these requirements, engineers have created specialized hardware such as Tensor Processing Units (TPUs) and neuromorphic chips that are optimized for these tasks. These are application specific architectures built for high task parallelism and low latency, highly efficient AI computations. This paper investigates how TPUs and neuromorphic chips fill in these gaps, outlines their designs and applications, and investigates their role in the advancement in the computer architectural design.

Main Body

-The Unique Demands of AI and ML Workloads

Matrix Multiplication Units Every TPU has a unique feature which sets it apart from its competitors and that is the speed at which it can perform matrix multiplications. Matrix multiplication is at the center of the training of neural networks, where the multi-dimensional array molecules represent the weights and activations.

To avoid the unnecessary, multi-purpose instructions, TPUs have specialized hardware that performs these computations. Simplified Control Logic TPUs cut down on control logic while AMD GPUs have a bit wider focus scope as TPUs use for a narrow range of usable control and only must focus on the computation parts of AI tasks. This design follows the above and increases throughput and energy efficiency. Generations of TPUs Systolic array architecture is one that TPUs make use of, where the flow of data is through the interconnected processing elements, and this helps to alleviate memory bottlenecks and improves performance.

TPU v1: Only able to carry out inference tasks TPU v1 was able to outperform the GPU of the time of 2016, by a factor of 15 to 30 times, with an amazing low power draw.

Final Project – Fall 2024

TPU v2 and iv3: Added support for training workloads and liquid cooling systems to cope with increased thermal loads.

TPU v4: The fourth and current generation is aimed at high scalability amid parallelism to increase efficiency for large scale AI tasks, like language model training with many billions of parameters.

Characteristics of AI and ML Workloads

Parallelism: AI models are typically trained over large matrices and tensors, with millions or even billions of parameters.

Neural networks involve various layers of computation, and thus the hardware must support high throughput with large-scale parallel processing capabilities.

Low Latency: Real-time AI applications such as autonomous driving or video analytics need immediate processing for a timely response.

Energy Efficiency: AI models can be computationally expensive, especially deep learning; consequently, they require a high amount of energy.

-Tensor Processing Units (TPUs)

Tensor Processing Units (TPUs), introduced by Google in 2016, are designed specifically to accelerate machine learning tasks. As application-specific integrated circuits (ASICs), TPUs focus on deep learning tasks, particularly tensor-based operations, which form the foundation of neural networks.

Key Features

Multiplication Units Every TPU has a unique feature which sets it apart from its

Final Project – Fall 2024

competitors and that is the speed at which it can perform matrix multiplications. Matrix multiplication is at the center of the training of neural networks, where the multi-dimensional array molecules represent the weights and activations. To avoid the unnecessary, multi-purpose instructions, TPUs have specialized hardware that performs these computations.

Simplified Control Logic TPUs cut down on control logic while AMD GPUs have a bit wider focus scope as TPUs use for a narrow range of usable control and only has to focus on the computation parts of AI tasks. This design follows the above and increases throughput and energy efficiency.

Generations of TPUs

Systolic array architecture is one that TPUs make use of, where the flow of data is through the interconnected processing elements, and this helps to alleviate memory bottlenecks and improves performance.

TPU v1: Only able to carry out inference tasks TPU v1 was able to outperform the GPU of the time period of 2016, by a factor of 15 to 30 times, with an amazing low power draw.

TPU v2 and iv3: Added support for training workloads and liquid cooling systems to cope with increased thermal loads.

TPU v4: The fourth and current generation is aimed at high scalability amid parallelism to increase efficiency for large scale AI tasks, like language model training with many billions of parameters.

Use Cases

Natural Language Processing (NLP): TPUs act as the nerve center of large scale NLP models including Google's BERT and LaMDA through which machines can understand and generate human languages proficiently.

Final Project – Fall 2024

Recommendation Systems: Organizations utilize TPUs for the purpose of preparing recommendation systems that assess user behavior patterns and recommend content that is most probably of their interest.

Healthcare Imaging: Medical diagnostics then become the possible use case of training machines that recognize images and detect the presence of abnormalities in X-rays and MRIs.

-Neuromorphic Chips

Neuromorphic computing is an innovative concept that refers to AI hardware design inspired by the neural connections in the human brain.

In contrast to conventional processors that consume data instantaneously, neuromorphic chips adopt the approach of parallel and event-triggered, just like neural networks.

Key Features

Spike-Based Processing: The computers making use of the new technology process information via spikes, or discrete events, instead of traditional continuous signals. Through this technique, energy efficient computation is achieved and thus mimicking the neuron communication process.

Massive Parallelism: Besides the existence of a neural network and other connectivity, neuromorphic semiconductors are also capable of acting as a parallel processing machine, thereby making them a valuable tool for processing complex sensory tasks in real-time just like vision and audio.

Learning Capabilities: Neuromorphic chips enable on-device learning, meaning systems can adjust to and even start performing better without needing outside help
HealthCare Imaging & AI.

Final Project – Fall 2024

Examples of Neuromorphic Chips

IBM TrueNorth: Already equipped with one million artificial neurons and 256 million synapses, TrueNorth has the power to operate cognitively thanks to the ultra-low power consumption technique.

Intel Loihi: Loihi is a combination of 128 neuromorphic cores that allow for real-time learning, thereby enabling it to adapt its behaviour to the environment and is perfect for robotics and adaptive systems.

Use Cases

Edge AI: With neuromorphic chips expected to excel at edge devices like drones and robotics, the chips will put a premium on efficient power usage and skillful decision-making in real time.

Sensory Processing: Designed for audio signal processing and computer vision, applications like these exploit the brain-like efficiency of neuromorphic chips.

-Comparison of TPUs and Neuromorphic Chips

Feature	TPUs	Neuromorphic Chips
Primary Focus	Tensor operations for deep learning	Brain-inspired computing
Power Efficiency	Moderate	High
Use Case Suitability	Cloud AI and data centers	Edge devices and real-time tasks
Learning Capability	Pre-trained models	On-device learning

Final Project – Fall 2024

-Impacts on Computer Architecture

Trend towards Heterogeneous Frameworks: The characters in the computer systems and electronic equipment are very many nowadays. Processing units ranging from CPUs, GPUs, TPUs, and a plethora of other special purpose hardware are utilized in various circumstances to facilitate increased handling of workloads. This tendency is focused on the system and workload-specific design of the hardware.

Power Consumption Considerations: The continuing growth and complexity of AI models makes energy efficiency an even more central interest. More niche AI hardware drives this evolution within which architecture is more favorably adapted.

Memory Advances: Such applications cause caretaking of the peculiarities of the memory technology development, for example, the introduction of high-bandwidth memory (HBM) to store large models in memory without read delays.

Interconnect: Starting with training very large language models like GPT-4 and moving on to supporting edge AIs in low power IoT devices, the hardware configuration in process has been undergoing changes.

Trend towards Heterogeneous Frameworks: The characters in the computer systems and electronic equipment are very many nowadays. Processing units ranging from CPUs, GPUs, TPUs, and a plethora of other special purpose hardware are utilized in various circumstances to facilitate increased handling of workloads. This tendency is focused on the system and workload-specific design of the hardware.

Power Consumption Considerations: The continuing growth and complexity of AI models makes energy efficiency an even more central interest. More niche AI hardware drives this evolution within which architecture is more favorably adapted.

Final Project – Fall 2024

Memory Advances: Such applications cause caretaking of the peculiarities of the memory technology development, for example, the introduction of high-bandwidth memory (HBM) to store large models in memory without read delays.

Interconnect: Starting with training very large language models like GPT-4 and moving on to supporting edge AIs in low power IoT devices, the hardware configuration in process has been undergoing changes.

Conclusion

Within the domain of objects that differ greatly to AI, such as AI-centric hardware advanced by Google such as Tensor Processing Units (TPUs) and biological supercomputer chips for bionic cyber-physical systems made from dry or wet neurons regeneratively grown chips, have become the benchmarks in the modern Computer space. These architectures unite aesthetic aspects of functional organizations which AI and AD together may ask, offering the best of both worlds for the aspired performance, while keeping the efficiency, and flexibility, at peak levels.

Cloud-based applications and experiments do not enjoy these advantages as the TPUs are more focused on the datacenter addressing their limitation.

There is an exceptional computer-generated prospect that the edge AI, will probably benefit most from this form of computing, and that is the use of neuromorphic chips.

References

IBM Research. (2016). TrueNorth: The Brain-Inspired Computing Chip. Retrieved from ibm.com.

Intel Corporation. (2021). Loihi Neuromorphic Computing. Retrieved from intel.com.

Final Project – Fall 2024

Jouppi, N. P., Young, C., Patil, N., Patterson, D., & Agrawal, G. (2017). In-datacenter performance analysis of a Tensor Processing Unit. *Proceedings of the 44th Annual International Symposium on Computer Architecture*.